

The generality of self-control

Chen, Kent
kentschen

Lee, Rachel
reychil

LeRoy, Benjamin
benjaminleroy

Liang, Jane
janewliang

Udagawa, Hiroto
hiroto-udagawa

November 9, 2015

Abstract

Self control is an interesting field of research with neurological facets that can be captured using the fMRI method. Relating a subjects ability to control themselves to an area of the brain or focusing on understanding what connects self control to neurological activity has been the focus of many studies. Our goal is to reproduce and extend the analysis for a Balloon Analogue Risk Task study found in *The Development and Generality of Self-Control* [?].

1 Introduction

This paper, *The Development and Generality of Self-Control* [1], and its associated fMRI studies are concerned with the relationship between impaired and normal self control as well as the similarities and differences across the brain relating to self-control. The entire paper explores multiple studies (of multiple study types), but we will just focus on the third and final study, which compares four different types of self control among healthy adults to see if they are related to each other. Very little relationship was found between these behavioral tasks, in contrast to the vast majority of existing literature, which argues for a unified notion of self-control, which is why we have decided to narrow the data analysis approaches to just the BART study. The BART study primarily focused on the inflation/deflation of a balloon that could pop; the fMRI scans show blood flow to the brain in an attempt to reveal control over risk-taking behavior while subjects participate in this particular study.

The rest of this report will go further in detail on the processes that we have tried to mimic from the original analysis of the data, but here is a brief overview of the work we have accomplished so far. We have looked into spatial smoothing of the data on the voxels of the brain. After realizing that there are many factors that could contribute to noise in the data, we decided it would be beneficial to look into smoothing modules in Python that could handle n-dimensional datasets to clarify which voxels are more important than other voxels. After obtaining convolved time courses for each subject, we turned to fitting simple and multiple linear regression models to each subject. Since examining the behavior of blood flow in the voxels over time was of such great interest, we also considered modeling the behavior as a time series using an autoregressive integrated moving average (ARIMA).

2 Data

The Balloon Analogue Risk Task (BART) study data consists of the average number of pumps for each balloon, how times the subject elected to explode the balloon, and how much money was earned across each trial for each respective participant. There are four different behavioral variables: the average number of pumps for each balloon, the average across each run, the number of exploded balloons, and the number of total trials. Within each variable, there are three model conditions: events for inflating the balloon (excluding the very last inflation of each trial), the last inflation before an explosion, and the event when the participant cashes out. There were 24 subjects with purportedly "clean" data.

3 Methods

3.1 Convolution

Our study has an event-related neurological stimulus, not block stimulus which was discussed in class. We predict the hemoglobin response (HR) related to this type of neurological stimulus, since the methods in class were much better for block stimulus. Moreover our stimulus (as is the case in event-related stimulus studies) wasn't evenly split. We developed a few other functions to take this into account, and will probably implement one more procedure before the end.

The three approaches we will have explored (the first two already completed) address modeling the HR from the neural stimulus. **(1)** We first implemented the standard block-stimulus paradigm related `np.convolve` function (which failed to take into account the non-trivial times of stimulus). `Np.convolve` utilizes FFT (Fast Fourier Transforms). **(2)** Next, we developed a function that doesn't utilize FFT, but allows for dis-continuous events of stimulus (discontinuous, being non-uniform stimulus timing). We used a continuous HR function to model the response (whereas `np.convolve` uses a discrete HR function). **(3)** Our next project (*which really should be included in the "Future plans" section*) is to utilize the strength of FFT by splitting up the time into 60 milliseconds intervals (a very small time metric), and then uses `np.convolve` and some rounding of the stimulus response occurrences to get another candidate for the Hemoglobin response related to the neurological stimulus. This procedure would use the strength of FFT, and wouldn't be hurt by assumptions of time since the HR model isn't super precise and there is a good amount of variance.

For **(2)** and **(3)** we have to scale back down to the image capturing time frame (2 seconds) [Figure 1].

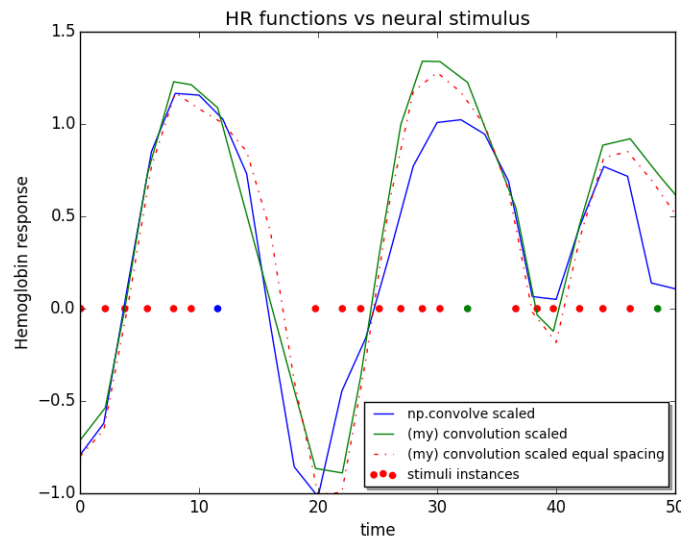


Figure 1: Different convolution function vs the Neural stimulus

3.2 Smoothing

Due to the random nature of human subjects and their movements, a certain extent of smoothing must be performed on the spatial dataset so that the noisy data can be cast off from the data that represent significant changes in blood flow in the brain. By doing so, researchers and anyone else investigating the data will be able to distinguish between non-brain scans versus actual brain scans. Each voxel of the brain is represented on a measure of blood flow intensity, and so a series of steps must be taken so that it is correctly convolved to most closely and accurately depict what was happening at a certain point in the brain at a certain time. After researching quite extensively, we decided to use a convolution involving a Gaussian kernel in order to smooth the three dimensional data. Originally, we were going to try and write a smoothing function from scratch, by implementing a rudimentary

rudimentary average-over-neighbors method. However, discussions with mentors lead us to the `scipy` module called `ndimage.filters` that has a function that performs a Gaussian filter on n-dimensional data. This was exactly what we needed so rather than reinventing the wheel, we will be smoothing the data with this module.

3.3 Regression

A simple and straightforward way to model the voxel time courses is to perform simple and multiple linear regressions. As a first attempt, we implemented and performed simple regression on a single subject's 4-dimensional array of voxels against the convolved time course, such that every voxel had an intercept and a coefficient corresponding to the convolved time course. However, examining the effects of the BART experiment conditions on voxel blood flow is also of interest. Thus, we turned to a more sophisticated multiple linear regression model that includes the conditions as dummy variable predictors. In either scenario, we created a design matrix X with the number of rows equal to the number of observed times and the number of columns is equal to the number of predictors. Then, using matrix algebra, we found the matrix of coefficients β by calculating $\beta = (XX^T)^{-1}X^TY$, where Y is the 4-dimensional array of the subject's voxels transformed into two-dimensional space. To do this, we essentially flattened out the first three dimensions (which indicate spatial positions) into a single dimension, while keeping the fourth dimension (time) the same. The resulting β was transformed into a three-dimensional array to maintain the spatial relationships of the voxels.

To consider the strength of the effects of these predictors, we looked at t-tests of the corresponding estimated coefficients for each voxel, as discussed under "*Hypothesis Testing*". Neither of these models, based on our current convolution methods, was very fruitful.

3.4 Hypothesis Testing

We created our simple linear regression model not necessarily for prediction but rather to understand the relationship between the voxels in a subject's brain and the convolved time course. In order to measure the strength of the relationship between these two measurements, we ran a hypothesis test on the coefficient of the linear regression model.

In our case, we ran a t-test on the β_1 coefficient of the model with the null hypothesis that $\beta_1 = 0$ and the alternative hypothesis that $\beta_1 \neq 0$. Since there is a linear model associated with each voxel in a subject's image, we have a single t-statistic associated to each voxel in our simple linear regression case. Once we had obtained the t-statistic, we compared this value across voxels in two ways. First, we simply compared this value across voxels in a subject. In this case, we took into account the sign of the t-statistic in our analysis. Second, we converted this t-statistic to a p-value, in which case the sign of the t-value will become irrelevant and we compared across voxels without taking into account this sign. We also developed functions to create t-statistics for multiple feature linear regression.

Now that we created a method to compare the voxels within a single subject, we next examined our results for the same voxels across subjects. Our initial step was to aggregate the t-statistic data between all subjects for each voxel. This allowed us to decrease the variability of the fit on each voxel and detect a more clear signal.

In order to do this, we ran the hypothesis test as stated above on all 24 subjects of the study. Then per voxel, we took the simple average of these values. An issue with our data was the presence of empty space that the scanner detected that is not directly part of the brain. In order to account for this, we took the mask data of the brain and cut out the parts of the images that were not relevant.

3.5 PCA

To perform Principal Components Analysis (PCA), we first considered the method outlined by the guide posted on the class website. Using matrix algebra, we were able to compute principal components by finding orthogonal projections (the shortest distance between points). which has the advantage of projecting the original data so that the predictors are independent.

After consulting with J.B. Poline in lecture, we took a new approach to PCA by implementing the Singular Value Decomposition (SVD) method for 4-dimensional data reshaped to 2-dimensional space. The SVD method is considerably more computationally efficient. With our large number predictors (the time intervals), the original 2-dimensional matrix may be too large to interpret or study because there

would be too many correlations to look at. To analyze the data in a more meaningful form, we were able to use PCA to reduce the number of important variables to a few, interpretable linear combinations of our dataset. Though the dimensions of the dataset did not change, we were able to make our dataset easier to explore and visualize. For example, with the help of eigenvectors and eigenvalues, PCA allowed us to focus on the components with the most explained variance.

3.6 Time Series

Cohen’s paper discusses analyzing the data with time series using FILM (FMRIBs Improved Linear Model). While we are not familiar with the FILM method, we did try modeling individual voxels in the framework of an autoregressive integrated moving average (ARIMA) process. We focused only on a single voxel from the first subject, but the method could easily be expanded to additional or aggregate voxels. Let $\{Y_t\}$ be a single volume’s value at time t and assume that the d th difference $W_t = \nabla^d Y_t$ is weakly stationary, defined to be when W_t has a constant mean function and autocovariance dependent only on lag k and not time t . Then we can try to model W_t as a linear combination of p autoregressive terms (or the number of most recent values to include) and q moving average terms (the number of lags to include for the white noise error terms):

$$W_t = \phi_1 W_{t-1} + \phi_2 W_{t-2} + \dots + \phi_p W_{t-p} + e_t - \theta_1 e_{t-1} - \theta_2 e_{t-2} - \dots - \theta_q e_{t-q}.$$

White noise is defined as a sequence of independent, identically distributed random variables. In order to fit an ARIMA process, the three orders p , d , and q must be first be specified, and the the associated coefficients estimated. We used a combination of visual inspection and quantitative methods to specify the ARIMA orders, and then used the maximum likelihood method to estimate parameters.

4 Results

4.1 Regression Results

To develop linear models, we looked at the HR from the neural response as a single feature and used multiple regression to take into account the 3 different types of stimulus (pump, explode, save) to see if the separation of these stimuli can better describe the response (specifically allowing for the prediction of different amplitudes for each type of stimulus’s HR function in linear regression). The predictive power of these models was not very good [Figure 2] [Figure 3].




Figure 2: Fitted vs Actual HR based on regressions




Figure 3: Plotting all predicted HR for conditions

As we also obtained beta values (coefficients) from the linear regression models, we looked at the 3-dimensional reports of the β values, a less rigorous analysis than hypothesis testing with t-statistics [Figure 4].

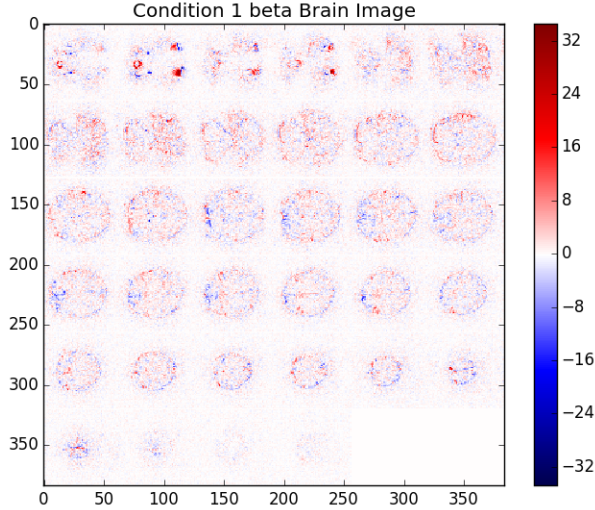


Figure 4: Beta values for condition 1, subject 001

In the future we will include more features to expand beyond information from the neural stimulus, as discussed in the *Future Work Section*.

4.2 Hypothesis Testing Results

We created our simple linear regression model not necessarily for prediction but rather to understand the relationship between the voxels in a subject's brain and the convolved time course. In order to measure the strength of the relationship between these two measurements, we ran a hypothesis test on the coefficient of the linear regression model.

In our case, we ran a t-test on the β_1 coefficient of the model with the null hypothesis that $\beta_1 = 0$




Figure 5: Ben sucks.

and the alternative hypothesis that $\beta \neq 0$. Since there is a linear model associated with each voxel in a subjects image, we have a single t-statistic associated to each voxel in our simple linear regression case. Once we had obtained the t-statistic, we compared this value across voxels in two ways. First, we simply compared this value across voxels in a subject. In this case, we took into account the sign of the t-statistic in our analysis. Second, we converted this t-statistic to a p-value, in which case the sign of the t-value will become irrelevant and we compared across voxels without taking into account this sign. We also developed functions to create t-statistics for multiple feature linear regression.

Now that we created a method to compare the voxels within a single subject, we next examined our results for the same voxels across subjects. Our initial step was to aggregate the t-statistic data between all subjects for each voxel. This allowed us to decrease the variability of the fit on each voxel and detect a more clear signal.

In order to do this, we ran the hypothesis test as stated above on all 24 subjects of the study. Then per voxel, we took the simple average of these values. An issue with our data was the presence of empty space that the scanner detected that is not directly part of the brain. In order to account for this, we took the mask data of the brain and cut out the parts of the images that were not relevant.

[Figure 5]

4.3 Time Series Results

We considered only a single voxel from the first subject. To specify the orders p , d , and q for an ARIMA process, we first checked for stationarity by considering the mean function and autocovariance. Visual inspection suggested that the mean function was nonconstant, and additional histograms and quantile-quantile plots suggested a slight skewness with a long right tail. We then corrected the skewness using a log transformation, deemed appropriate by the Box-Cox method. The log-transformed data still exhibited a non-constant mean, so we considered using the first difference (i.e. $W_t = Y_t - Y_{t-1}$). The first difference appeared to be much more reasonably stationary. So, we took $d = 1$.

Having specified the order for d , we turned to the problem of specifying p and q . We used a combination of visually inspecting the autocorrelation and partial autocorrelation plots of the first difference, and looking at the AIC and BIC computed from a grid of possible models. The latter method suggested specifying $p = 1$ and $q = 1$ (based on either the AIC or the BIC), which was also supported by the visual inspections.

We estimated the parameters for an ARIMA(1,1,1) model using the exact maximum likelihood estimator via Kalman filter. The residuals appear to be normally distributed, and its autocorrelation and partial autocorrelation plots also do not raise any red flags. Furthermore, when visually comparing the fitted time series to the true observed data, the ARIMA process seems to approximate the observed

data much better than the linear regression models. While more work, such as developing more robust methods to assess fit and considering the problem of modeling multiple voxels across multiple subjects, is clearly needed, time series analysis presents a promising direction for further investigation.

In particular, we may be able to forecast future observations based on previous ones. As an example we modeled an ARIMA(1,1,1) process based on the first half of the observations for a single voxel. This process was then used to forecast the second half of the observations. A comparison between the true observations and the forecasted predictions is shown in [Figure 6]. While the forecasted observations look reasonable for approximating the true values, more quantitative metrics for assessing performance need to be implemented.




Figure 6: Forecasting the second half of observations based on the first half.

5 Discussion

5.1 Discussion of Results

While very much still a work in progress, our analysis thus far includes both data processing and methods for modeling voxel time courses. Prior to doing any serious analysis, we had to smooth the data spatially for each subject and also generate a reasonable convolved time course based on event-related neurological stimulus with non-constant intervals. A simple, but nevertheless important model to consider is linear regression. We implemented both simple and multiple regression models at the individual subject level, and performed hypothesis testing on the resulting coefficients for each voxel. Finally, we explored some additional methods that may be of interest or use moving forward, including principal components analysis and modeling individual volumes as a time series.

5.2 Discussion of Future Work

Extending on our work with multiple linear regression we will explore modeling more of the noise in our data, with linear drift and discrete cosine transforms for general trends (similar to Fourier series as extra features).

We will also be looking into (as long as this isn't corrected with pre-processing from mentors), realignment of scans to correct for the time it takes to scan each voxel compared to the start of the scan.

One major concern for hypothesis testing and working with the estimated coefficients from linear modelling is the issue of multiple comparisons. This may be as simple as utilizing a Bonferroni correction, but we can also consider permutation tests or more sophisticated models (Benjamini-Hoffberg, etc).

More quantitative and robust indicators for validating time series models should be implemented. One possibility is to simulate a null process by permuting the phases of the voxels time course after performing a Fourier transform, and then transforming the permuted data back to the original space. This way, the data is permuted but maintains the original correlated structure. We can then fit the ARIMA model to the permuted process and examine how much, if at all, the ARIMA process fitted to the observed data makes improvements over the null case. Generating confidence intervals for the parameter estimates, expanding the procedure across multiple voxels or subjects, and forecasting future observations may also be of interest.

References

- [1] J. R. COHEN, *The development and generality of self-control*, ProQuest, (2009), p. 164.